

Notes 8: Model Selection

Nathaniel E. Helwig

Department of Statistics
University of Illinois at Urbana-Champaign

Stat 420: Methods of Applied Statistics
Section N1U/N1G – Spring 2014

Outline of Notes

1) Model Selection Problem:

- Overview
- Wine Example
- Selection Summary

2) P-Value Model Selection:

- Backwards Elimination
- Forward Selection
- Stepwise Selection

3) Intelligent Model Selection:

- Adjusted R^2
- Information Criteria
- Prediction Based

4) State Data Example:

- Data Description
- P-Value Selection
- Intelligent Selection

The Problem: Which Variables??

The problem of model selection asks the question: which variables should be included in a multiple regression model?

We do not want to include too many predictors.

- Problem of over-fitting data
- Solution may not cross-validate

We do not want to include too few predictors.

- Miss important relationships in data
- Misinterpret relationships in data

All Possible Models

We need to consider ALL possible models that could be formed.

If we have p predictors, then (according to binomial theorem) there are

$$\sum_{j=1}^p \binom{p}{j} = 2^p$$

possible models we could choose.

Wine Example: Overview

Using Wine Quality data set from UCI Machine Learning:
<http://archive.ics.uci.edu/ml/datasets/Wine+Quality>

Predict wine quality from 11 possible physicochemical predictors:

- quality (range 0-10)
- fixed acidity (g/dm^3)
- volatile acidity (g/dm^3)
- citric acid (g/dm^3)
- residual sugar (g/dm^3)
- chlorides (g/dm^3)
- free sulfur dioxide (mg/dm^3)
- total sulfur dioxide (mg/dm^2)
- density (g/cm^3)
- pH
- sulphates (g/dm^3)
- alcohol (% ABV)

Wine Example: Look at Data

```
> wines=read.table("/Users/Nate/Desktop/winequality-red.csv",
                  sep=";",header=TRUE)
> wines[1:3,]
  fixed.acidity volatile.acidity citric.acid residual.sugar
1          7.4           0.70         0.00           1.9
2          7.8           0.88         0.00           2.6
3          7.8           0.76         0.04           2.3
 chlorides free.sulfur.dioxide total.sulfur.dioxide density
1      0.076              11              34 0.9978
2      0.098              25              67 0.9968
3      0.092              15              54 0.9970
   pH sulphates alcohol quality
1 3.51      0.56      9.4      5
2 3.20      0.68      9.8      5
3 3.26      0.65      9.8      5
```

Wine Example: Predicting Red Wine Quality

```
> wmod=lm(quality~.,data=wines)
> summary(wmod)
```

```
Call:
lm(formula = quality ~ ., data = wines)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.68911	-0.36652	-0.04699	0.45202	2.02498

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.197e+01	2.119e+01	1.036	0.3002
fixed.acidity	2.499e-02	2.595e-02	0.963	0.3357
volatile.acidity	-1.084e+00	1.211e-01	-8.948	< 2e-16 ***
citric.acid	-1.826e-01	1.472e-01	-1.240	0.2150
residual.sugar	1.633e-02	1.500e-02	1.089	0.2765
chlorides	-1.874e+00	4.193e-01	-4.470	8.37e-06 ***
free.sulfur.dioxide	4.361e-03	2.171e-03	2.009	0.0447 *
total.sulfur.dioxide	-3.265e-03	7.287e-04	-4.480	8.00e-06 ***
density	-1.788e+01	2.163e+01	-0.827	0.4086
pH	-4.137e-01	1.916e-01	-2.159	0.0310 *
sulphates	9.163e-01	1.143e-01	8.014	2.13e-15 ***
alcohol	2.762e-01	2.648e-02	10.429	< 2e-16 ***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 0.648 on 1587 degrees of freedom
Multiple R-squared:  0.3606, Adjusted R-squared:  0.3561
F-statistic: 81.35 on 11 and 1587 DF,  p-value: < 2.2e-16
```

Wine Example: Problems with Model

Model on previous slide is very complicated, so interpretation of the coefficients is problematic.

- Remember conditional interpretation of b_j parameters

Model on previous slide has included many effects which may not be needed to predict wine `quality`.

- Model is not likely to cross-validate well.

How do we select which of the physicochemical predictors are needed to predict wine `quality`??

The Solution: Model Selection Strategies

We can use different statistical model selection strategies to choose which predictors to include.

There are a variety of strategies we can use:

- P-value based methods (NOT good)
- Adjusted R^2 (better)
- Information criteria (best)
- Prediction/cross-validation (best)

Overview of P-Value Model Selection

P-value model selection strategies choose which terms to include based on the significance of the terms (i.e., p-values of F tests).

There are three popular p-value based selection strategies:

- Backwards elimination
- Forward selection
- Stepwise selection

There is no guarantee that these selection strategies will produce a reasonable (or the same) model!

Backwards Elimination Algorithm

Given a threshold α^* , backwards elimination algorithm is:

- 1 Begin with all possible predictors in model
- 2 Remove predictor with largest p-value above α^*
- 3 Refit and repeat step 2 until all p-values below α^*

Note that α^* doesn't have to be the magical 0.05; typically set α^* larger (e.g., 0.10 or 0.15) if ultimately interested in prediction.

- Do not want to miss important predictors

Backwards Elimination: R Function

We can write our own R function to perform backwards elimination:

```
bslm<-function(y,X,alpha=0.1,maxstep=NA,prindrop=TRUE,stepc=0L){
  X=data.frame(X);      xnames=names(X)
  mymod=lm(y~.,data=X)
  if(is.na(maxstep)){maxstep=length(xnames)+stepc}
  notsig=TRUE;          nstep=0L+stepc
  while(notsig && nstep<maxstep){
    nstep=nstep+1L
    mydrop=drop1(mymod,test="F")
    pvals=mydrop[,6];    mxidx=which.max(pvals)
    if(pvals[mxidx]>alpha){
      dropname=rownames(mydrop)[mxidx]
      myform=as.formula(paste("~.",dropname))
      mymod=update(mymod,myform)
      if(prindrop){print(paste("Step",nstep,":","Dropping",dropname))}
    } else{notsig=FALSE}
  }
  mymod
}
```

Backwards Elimination: Wine Example

```
> bswmod=bslm(wines$quality,wines[,-12])
[1] "Step 1 : Dropping density"
[1] "Step 2 : Dropping fixed.acidity"
[1] "Step 3 : Dropping residual.sugar"
[1] "Step 4 : Dropping citric.acid"
> summary(bswmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom
Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567
F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

Note that we dropped

- ❶ density
- ❷ fixed.acidity
- ❸ residual.sugar
- ❹ citric.acid

Forward Selection Algorithm

Given a threshold α^* , forward selection algorithm is:

- 1 Begin with no predictors in model
- 2 Add predictor with smallest p-value below α^*
- 3 Refit and repeat step 2 until no new p-values below α^*

Note that α^* doesn't have to be the magical 0.05; typically set α^* larger (e.g., 0.10 or 0.15) if ultimately interested in prediction.

- Do not want to miss important predictors

Forward Selection in R

We can write our own R function to perform forward selection:

```

fslm<-function(y,X,alpha=0.1,maxstep=NA,printadd=TRUE,stepc=0L,indx=NULL){
  X=data.frame(X);      xnames=names(X)
  if(is.null(indx)){
    mymod=lm(y~1);  anames=NULL
  } else {
    anames=xnames[indx]; xnames=xnames[-indx]
    myform=as.formula(paste("y",paste(anames,collapse="+"),sep=~"))
    mymod=lm(myform,data=X)
  }
  if(is.na(maxstep)){maxstep=length(xnames)+stepc}
  notsig=TRUE;  nstep=0L+stepc
  while(notsig && nstep<maxstep){
    nstep=nstep+1L
    myform=as.formula(paste("~",paste(c(anames,xnames),collapse="+")))
    myadd=add1(mymod,scope=myform,x=model.matrix(myform,data=X),test="F")
    pvals=myadd[,6];  mnidx=which.min(pvals)
    if(pvals[mnidx]<alpha){
      anames=c(anames,rownames(myadd)[mnidx])
      addname=rownames(myadd)[mnidx]
      myform=as.formula(paste(".~."+addname))
      mymod=update(mymod,myform,data=X)
      if(printadd){print(paste("Step",nstep,":", "Adding", addname))}
      xnames=xnames[-(mnidx-1L)]
    } else{notsig=FALSE}
  }
  mymod
}

```

Forward Selection: Wine Example

```
> fswmod=fslm(wines$quality,wines[,-12])
[1] "Step 1 : Adding alcohol"
[1] "Step 2 : Adding volatile.acidity"
[1] "Step 3 : Adding sulphates"
[1] "Step 4 : Adding total.sulfur.dioxide"
[1] "Step 5 : Adding chlorides"
[1] "Step 6 : Adding pH"
[1] "Step 7 : Adding free.sulfur.dioxide"
> summary(fswmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

Note we did NOT add

- fixed.acidity
- citric.acid
- residual.sugar
- density

Stepwise Selection Algorithm

Given thresholds α_F^* and α_B^* , stepwise selection algorithm is:

- 1 Begin with no predictors in model
- 2 Forward step: add predictor with smallest p-value below α_F^*
- 3 Backward step: remove predictor with largest p-value above α_B^*
- 4 Repeat steps 2–3 until convergence (or max steps reached)

Note that α_F^* and α_B^* do not have to be the magical 0.05; typically set α^* larger (e.g., 0.10 or 0.15) if ultimately interested in prediction.

- α_F^* and α_B^* do not have to be equal

Stepwise Selection in R

We can write our own R function to perform stepwise selection:

```
swlm<-function(y,X,alpha=c(0.1,0.1),maxstep=20,printad=TRUE){
  X=data.frame(X);      xnames=names(X)
  notsig=TRUE;  nstep=0L;  inidx=oldidx=NULL
  while(notsig && nstep<maxstep){
    nstep=nstep+1L
    # forward step
    mymod=fslm(y,X,alpha[1],nstep,printad,(nstep-1L),inidx)
    innames=attr(mymod$terms,"term.labels")
    inidx=match(innames,xnames)
    # backwards step
    if(nstep>1L){
      Xin=data.frame(X[,inidx]);  names(Xin)=innames
      mymod=bslm(y,Xin,alpha[2],nstep,printad,(nstep-1L))
      innames=attr(mymod$terms,"term.labels")
      inidx=match(innames,xnames)
    }
    # check convergence
    if(length(inidx)==length(oldidx) && any(inidx!=oldidx)==FALSE){notsig=FALSE}
    oldidx=inidx
  }
  mymod
}
```

Stepwise Selection: Wine Example

```
> stwmod=swlm(wines$quality,wines[, -12])
[1] "Step 1 : Adding alcohol"
[1] "Step 2 : Adding volatile.acidity"
[1] "Step 3 : Adding sulphates"
[1] "Step 4 : Adding total.sulfur.dioxide"
[1] "Step 5 : Adding chlorides"
[1] "Step 6 : Adding pH"
[1] "Step 7 : Adding free.sulfur.dioxide"
> summary(stwmod) # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

Note we did not include

- fixed.acidity
- citric.acid
- residual.sugar
- density

Selection Summary: Wine Example

Bad example where all three methods choose same model.

- Note that selected model is NOT easily interpretable
- Note that selected model is NOT theoretically motivated

Although we have “systematically selected” this model, there is no guarantee that this is a good (i.e., valid and useful) model.

It would be best to use some theory to guide model building.

Selection Summary: Toy Example

```
> set.seed(1120)
> x1=rnorm(100)
> x2=rnorm(100)
> x3=(x1+x2)/2+rnorm(100,sd=.1)
> X=as.data.frame(cbind(x1,x2,x3))
> y=x1+x2+rnorm(100)

> bstest=bslm(y,X)      # picks correct model
[1] "Step 1 : Dropping x3"
> fstest=fslm(y,X)      # picks wrong model
[1] "Step 1 : Adding x3"
> swtest=swlm(y,X)      # picks wrong model
[1] "Step 1 : Adding x3"
```

Coefficient of Multiple Determination (revisited)

Consider the MLR model $y_i = b_0 + \sum_{j=1}^p b_j x_{ij} + e_i$ with $e_i \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$.

Remember: the *coefficient of multiple determination* is defined as

$$\begin{aligned} R^2 &= \frac{SSR}{SST} \\ &= 1 - \frac{SSE}{SST} \end{aligned}$$

and gives the amount of variation in y_i that is explained by the linear relationships with x_{i1}, \dots, x_{ip} .

Adjusted R^2 (revisited)

Including more predictors in a MLR model can artificially inflate R^2 :

- Capitalizing on spurious effects present in noisy data
- Phenomenon of *over-fitting* the data

The *adjusted* R^2 is a relative measure of fit:

$$\begin{aligned} R_a^2 &= 1 - \frac{SSE/df_E}{SST/df_T} \\ &= 1 - \frac{\hat{\sigma}^2}{s_Y^2} \end{aligned}$$

where $s_Y^2 = \frac{\sum_{i=1}^n (y_i - \bar{y})^2}{n-1}$ is the sample estimate of the variance of Y .

Adjusted R^2 for Model Selection

If p is not too large, could calculate R_a^2 for all 2^p possible models.

- Pick model with largest R_a^2 .

Implemented in `leaps` function (`leaps` package).

- Branch-and-bound search through all possible subsets
- Use `method="adjr2"` option to select via adjusted R^2

Adjusted R^2 Selection: Wine Example

```
> X=wines[,-12]
> arsqmod=leaps(x=X,y=wines$quality,method="adjr2")
> widx=which.max(arsqmod$adjr2)
> xidx=(1:ncol(X))[arsqmod$which[widx,]]
> Xin=data.frame(X[,xidx])
> arsqmod=lm(wines$quality~.,data=Xin)
```

Adjusted R^2 Selection: Wine Example (continued)

```
> summary(arsqmod)           # I deleted some output
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.6680876  0.4608410  10.129 < 2e-16 ***
volatile.acidity -1.0736123  0.1159362  -9.260 < 2e-16 ***
citric.acid    -0.1295444  0.1217717  -1.064  0.2876
chlorides     -1.9494185  0.4026906  -4.841 1.42e-06 ***
free.sulfur.dioxide  0.0047601  0.0021463   2.218  0.0267 *
total.sulfur.dioxide -0.0033658  0.0006954  -4.840 1.42e-06 ***
pH            -0.5491501  0.1331350  -4.125 3.90e-05 ***
sulphates      0.8914283  0.1102122   8.088 1.19e-15 ***
alcohol        0.2928780  0.0171280  17.099 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1590 degrees of freedom
Multiple R-squared:  0.3599, Adjusted R-squared:  0.3567
F-statistic: 111.8 on 8 and 1590 DF,  p-value: < 2.2e-16

> summary(stwmod)$adj
[1] 0.3566527
> summary(arsqmod)$adj
[1] 0.356706
```

Likelihood Function (revisited)

Remember that $(\mathbf{y}|\mathbf{X}) \sim N(\mathbf{X}\mathbf{b}, \sigma^2\mathbf{I}_n)$, which implies that \mathbf{y} has pdf

$$f(\mathbf{y}|\mathbf{X}, \mathbf{b}, \sigma^2) = (2\pi)^{-n/2}(\sigma^2)^{-n/2} e^{-\frac{1}{2\sigma^2}(\mathbf{y}-\mathbf{X}\mathbf{b})'(\mathbf{y}-\mathbf{X}\mathbf{b})}$$

As a result, the log-likelihood of (\mathbf{b}, σ^2) given (\mathbf{y}, \mathbf{X}) is

$$\ln\{L(\mathbf{b}, \sigma^2|\mathbf{y}, \mathbf{X})\} = -\frac{n}{2}\ln(2\pi) - \frac{n}{2}\ln(\sigma^2) - \frac{1}{2\sigma^2}(\mathbf{y} - \mathbf{X}\mathbf{b})'(\mathbf{y} - \mathbf{X}\mathbf{b})$$

Maximized Likelihood Functions

Remember that the MLEs of \mathbf{b} and σ^2 are

$$\begin{aligned}\hat{\mathbf{b}} &= (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y} \\ \tilde{\sigma}^2 &= SSE/n\end{aligned}$$

where $SSE = (\mathbf{y} - \mathbf{X}\hat{\mathbf{b}})'(\mathbf{y} - \mathbf{X}\hat{\mathbf{b}})$ is the sum-of-squared errors.

As a result, the *maximized* log-likelihood of (\mathbf{b}, σ^2) given (\mathbf{y}, \mathbf{X}) is

$$\begin{aligned}\ln\{L(\hat{\mathbf{b}}, \tilde{\sigma}^2|\mathbf{y}, \mathbf{X})\} &= -\frac{n}{2}\ln(2\pi) - \frac{n}{2}\ln(\tilde{\sigma}^2) - \frac{1}{2\tilde{\sigma}^2}(\mathbf{y} - \mathbf{X}\hat{\mathbf{b}})'(\mathbf{y} - \mathbf{X}\hat{\mathbf{b}}) \\ &= -\frac{n}{2}\ln(2\pi) - \frac{n}{2}\ln(\tilde{\sigma}^2) - \frac{n}{2}\end{aligned}$$

Likelihoods and Information Criteria

Information criteria define model fit using maximized likelihoods that are penalized according to model complexity.

Defining $\hat{\mathcal{L}} = \ln\{L(\hat{\mathbf{b}}, \hat{\sigma}^2 | \mathbf{y}, \mathbf{X})\}$, Akaike's (1974) AIC is defined as

$$AIC = -2\hat{\mathcal{L}} + 2q$$

where q is number of parameters; note that AIC stands for *an information criterion*, but people typically refer to it as Akaike's.

The Bayesian Information Criterion (BIC; Schwarz, 1978) is

$$BIC = -2\hat{\mathcal{L}} + \ln(n)q$$

Information Criteria in Regression

Using the definition $\hat{\mathcal{L}} = -\frac{n}{2} \ln(2\pi) - \frac{n}{2} \ln(\tilde{\sigma}^2) - \frac{n}{2}$, we have that

$$AIC = n + n \ln(2\pi) + n \ln(\tilde{\sigma}^2) + 2c$$

$$BIC = n + n \ln(2\pi) + n \ln(\tilde{\sigma}^2) + \ln(n)c$$

where $c = p + 1$ is the number of columns of the model design matrix.

In some cases the constant $n + n \ln(2\pi)$ is dropped, such as

$$AIC = n \ln(\tilde{\sigma}^2) + 2c$$

$$BIC = n \ln(\tilde{\sigma}^2) + \ln(n)c$$

Information Criteria and Model Selection

AIC and BIC are theoretical optimal criteria for model selection.

- Smaller AIC (or BIC) means better model.
- $AIC < BIC$ whenever $n \geq 8 \implies AIC$ tends to pick larger models

AIC is optimal model selection criterion if trying to find model that best describes data among possible candidate models

- True model is unknown and not one of candidate models

BIC is optimal model selection criterion if trying to find true model among possible candidate models

- True model is one of candidate models

AIC and BIC Model Selection in R

You can perform AIC and BIC model selection using `step` function.

- Default use performs stepwise AIC selection
(`direction="both"` and `k=2`)
- Use `direction="backward"` or `direction="forward"` to change selection algorithm
- Set `k=log(n)` to perform BIC selection

Wine Example: Backward AIC Selection

```
> wmod=lm(quality~.,data=wines)
> bwine=step(wmod,direction="backward",trace=0)
> summary(bwine)
```

```
Call:
lm(formula = quality ~ volatile.acidity + chlorides + free.sulfur.dioxide +
    total.sulfur.dioxide + pH + sulphates + alcohol, data = wines)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.68918	-0.36757	-0.04653	0.46081	2.02954

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16	***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16	***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07	***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017	*
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07	***
pH	-0.4826614	0.1175581	-4.106	4.23e-05	***
sulphates	0.8826651	0.1099084	8.031	1.86e-15	***
alcohol	0.2893028	0.0167958	17.225	< 2e-16	***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

Wine Example: Forward AIC Selection

```
> wnames=names(wines)[-12]
> wmod=lm(quality~1,data=wines)
> wform=as.formula(paste("quality~",paste(wnames,collapse="+")))
> fwine=step(wmod,scope=wform,direction="forward",trace=0)
> summary(fwine)
```

Call:

```
lm(formula = quality ~ alcohol + volatile.acidity + sulphates +
    total.sulfur.dioxide + chlorides + pH + free.sulfur.dioxide,
    data = wines)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.68918	-0.36757	-0.04653	0.46081	2.02954

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

Wine Example: Stepwise AIC Selection

```
> wnames=names(wines)[-12]
> wmod=lm(quality~1,data=wines)
> wform=as.formula(paste("quality~",paste(wnames,collapse="+")))
> swine=step(wmod,scope=wform,trace=0)
> summary(swine)
```

Call:

```
lm(formula = quality ~ alcohol + volatile.acidity + sulphates +
    total.sulfur.dioxide + chlorides + pH + free.sulfur.dioxide,
    data = wines)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.68918	-0.36757	-0.04653	0.46081	2.02954

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

Wine Example: Backward BIC Selection

```
> wmod=lm(quality~.,data=wines)
> bwine=step(wmod,direction="backward",k=log(1599),trace=0)
> summary(bwine)
```

```
Call:
lm(formula = quality ~ volatile.acidity + chlorides + total.sulfur.dioxide +
    pH + sulphates + alcohol, data = wines)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.60575	-0.35883	-0.04806	0.46079	1.95643

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.2957316	0.3995603	10.751	< 2e-16 ***
volatile.acidity	-1.0381945	0.1004270	-10.338	< 2e-16 ***
chlorides	-2.0022839	0.3980757	-5.030	5.46e-07 ***
total.sulfur.dioxide	-0.0023721	0.0005064	-4.684	3.05e-06 ***
pH	-0.4351830	0.1160368	-3.750	0.000183 ***
sulphates	0.8886802	0.1100419	8.076	1.31e-15 ***
alcohol	0.2906738	0.0168108	17.291	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6487 on 1592 degrees of freedom
 Multiple R-squared: 0.3572, Adjusted R-squared: 0.3548
 F-statistic: 147.4 on 6 and 1592 DF, p-value: < 2.2e-16

Wine Example: Forward BIC Selection

```
> wnames=names(wines)[-12]
> wmod=lm(quality~1,data=wines)
> wform=as.formula(paste("quality~",paste(wnames,collapse="+")))
> fwine=step(wmod,scope=wform,direction="forward",k=log(1599),trace=0)
> summary(fwine)
```

Call:

```
lm(formula = quality ~ alcohol + volatile.acidity + sulphates +
    total.sulfur.dioxide + chlorides + pH, data = wines)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.60575	-0.35883	-0.04806	0.46079	1.95643

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.2957316	0.3995603	10.751	< 2e-16 ***
alcohol	0.2906738	0.0168108	17.291	< 2e-16 ***
volatile.acidity	-1.0381945	0.1004270	-10.338	< 2e-16 ***
sulphates	0.8886802	0.1100419	8.076	1.31e-15 ***
total.sulfur.dioxide	-0.0023721	0.0005064	-4.684	3.05e-06 ***
chlorides	-2.0022839	0.3980757	-5.030	5.46e-07 ***
pH	-0.4351830	0.1160368	-3.750	0.000183 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6487 on 1592 degrees of freedom

Multiple R-squared: 0.3572, Adjusted R-squared: 0.3548

F-statistic: 147.4 on 6 and 1592 DF, p-value: < 2.2e-16

Wine Example: Stepwise BIC Selection

```
> wnames=names(wines)[-12]
> wmod=lm(quality~1,data=wines)
> wform=as.formula(paste("quality~",paste(wnames,collapse="+")))
> swine=step(wmod,scope=wform,k=log(1599),trace=0)
> summary(swine)
```

Call:

```
lm(formula = quality ~ alcohol + volatile.acidity + sulphates +
    total.sulfur.dioxide + chlorides + pH, data = wines)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.60575	-0.35883	-0.04806	0.46079	1.95643

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.2957316	0.3995603	10.751	< 2e-16 ***
alcohol	0.2906738	0.0168108	17.291	< 2e-16 ***
volatile.acidity	-1.0381945	0.1004270	-10.338	< 2e-16 ***
sulphates	0.8886802	0.1100419	8.076	1.31e-15 ***
total.sulfur.dioxide	-0.0023721	0.0005064	-4.684	3.05e-06 ***
chlorides	-2.0022839	0.3980757	-5.030	5.46e-07 ***
pH	-0.4351830	0.1160368	-3.750	0.000183 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6487 on 1592 degrees of freedom

Multiple R-squared: 0.3572, Adjusted R-squared: 0.3548

F-statistic: 147.4 on 6 and 1592 DF, p-value: < 2.2e-16

Toy Example: AIC Selection

```

> set.seed(1120)
> x1=rnorm(100)
> x2=rnorm(100)
> x3=(x1+x2)/2+rnorm(100,sd=.1)
> X=as.data.frame(cbind(x1,x2,x3))
> y=x1+x2+rnorm(100)

> mymod=lm(y~.,data=X)
> amod=step(mymod,trace=0)
> amod$coef
      (Intercept)           x1           x2
-0.03036068    1.00370415    1.00912964

> mymod=lm(y~.,data=X)
> bmod=step(mymod,direction="backward",trace=0)
> bmod$coef
      (Intercept)           x1           x2
-0.03036068    1.00370415    1.00912964

> mymod=lm(y~1,data=X)
> fmod=step(mymod,y~x1+x2+x3,direction="forward",trace=0)
> fmod$coef
      (Intercept)           x3
-0.03617019    1.92491439

> mymod=lm(y~1,data=X)
> smod=step(mymod,y~x1+x2+x3,trace=0)
> smod$coef
      (Intercept)           x3
-0.03617019    1.92491439

```

Toy Example: BIC Selection

```

> set.seed(1120)
> x1=rnorm(100)
> x2=rnorm(100)
> x3=(x1+x2)/2+rnorm(100,sd=.1)
> X=as.data.frame(cbind(x1,x2,x3))
> y=x1+x2+rnorm(100)

> mymod=lm(y~.,data=X)
> amod=step(mymod,k=log(100),trace=0)
> amod$coef
      (Intercept)           x1           x2
-0.03036068    1.00370415    1.00912964

> mymod=lm(y~.,data=X)
> bmod=step(mymod,direction="backward",k=log(100),trace=0)
> bmod$coef
      (Intercept)           x1           x2
-0.03036068    1.00370415    1.00912964

> mymod=lm(y~1,data=X)
> fmod=step(mymod,y~x1+x2+x3,direction="forward",k=log(100),trace=0)
> fmod$coef
      (Intercept)           x3
-0.03617019    1.92491439

> mymod=lm(y~1,data=X)
> smod=step(mymod,y~x1+x2+x3,k=log(100),trace=0)
> smod$coef
      (Intercept)           x3
-0.03617019    1.92491439

```


Prediction and Model Selection

If we are ultimately interested in prediction, we can use prediction-based criteria to select our model.

Idea: minimize prediction SSE (instead of SSE for given data).

Most implementations do exhaustive (or branch-and-bound) searches, but you could use these criterion in a stepwise fashion too.

Mallow's C_p

Consider the model $\mathbf{y} = \mathbf{X}\mathbf{b} + \mathbf{e}$ where \mathbf{X} is $n \times m$ and $\mathbf{e} \sim N(\mathbf{0}, \sigma^2 \mathbf{I})$.

If we want to estimate the mean-squared prediction error (MSPE)

$$\frac{1}{\sigma^2} \sum_{i=1}^n E \left\{ [\hat{y}_i - E(y_i | \mathbf{x}_i)]^2 \right\}$$

we can use Mallow's (1973) C_p

$$C_p = \frac{SSE_p}{\hat{\sigma}^2} - n + 2p$$

where

- SSE_p is the SSE with $p < m$ columns of \mathbf{X} used in fit
- $\hat{\sigma}^2 = SSE/(n - m)$ is the MSE of full model

Mallow's C_p in R

Implemented in `leaps` function (`leaps` package).

- Branch-and-bound search through all possible subsets
- Use default `method="Cp"` option to select via Mallow's C_p

We could also use the `drop1` and `add1` functions

- These were called within `bslm`, `fslm`, `swlm`, and `step`
- Use `scale=` input to get Mallow's C_p

Mallow's C_p : Wine Example

```
> X=wines[,-12]
> cpmo=leaps(x=X,y=wines$quality,method="Cp")
> widx=which.min(cpmo$Cp)
> xidx=(1:ncol(X))[cpmo$which[widx,]]
> Xin=data.frame(X[,xidx])
> cpmo=lm(wines$quality~.,data=Xin)
> summary(cpmo)
```

Call:

```
lm(formula = wines$quality ~ ., data = Xin)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.68918	-0.36757	-0.04653	0.46081	2.02954

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16a

Mallow's C_p : Wine Example (continued)

```
> wmod=lm(quality~.,data=wines)
> sigmasq=summary(wmod)$sigma^2
> drop1(wmod,scale=sigmasq)
```

Single term deletions

Model:

```
quality ~ fixed.acidity + volatile.acidity + citric.acid + residual.sugar +
  chlorides + free.sulfur.dioxide + total.sulfur.dioxide +
  density + pH + sulphates + alcohol
```

scale: 0.4199185

	Df	Sum of Sq	RSS	Cp
<none>			666.41	12.000
fixed.acidity	1	0.389	666.80	10.928
volatile.acidity	1	33.620	700.03	90.063
citric.acid	1	0.646	667.06	11.539
residual.sugar	1	0.498	666.91	11.185
chlorides	1	8.391	674.80	29.982
free.sulfur.dioxide	1	1.694	668.10	14.035
total.sulfur.dioxide	1	8.427	674.84	30.069
density	1	0.287	666.70	10.683
pH	1	1.957	668.37	14.661
sulphates	1	26.971	693.38	74.229
alcohol	1	45.672	712.08	118.764

Drop density because it has the smallest C_p score below $p_0 = 12$.

Predicted Residual Sum-of-Squares (PRESS)

The Predicted Residual Sum-of-Squares (PRESS) statistic is

$$\text{PRESS} = \sum_{i=1}^n (y_i - \hat{y}_{[-i]})^2 = \sum_{i=1}^n \left(\frac{\hat{e}_i}{1 - h_{ii}} \right)^2$$

where

- $\hat{y}_{[-i]} = \mathbf{x}_i \hat{\mathbf{b}}_{[-i]}$ and $\hat{\mathbf{b}}_{[-i]}$ is estimate of \mathbf{b} without i -th observation
- \hat{e}_i is i -th estimated residual from full model
- h_{ii} is i -th leverage score from full model

PRESS Statistic in R

```
getpress<-function(indx,y,Xmat){  
  ivec=(1:ncol(X))[indx]  
  mymod=lm(y~.,data=data.frame(Xmat[,ivec]))  
  sum((mymod$residuals/(1-hatvalues(mymod)))^2)  
}
```

```
presslm<-function(y,X){  
  X=data.frame(X); np=ncol(X)  
  xlist=vector("list",np)  
  for(j in 1:np){xlist[[j]]=c(TRUE,FALSE)}  
  xall=expand.grid(xlist); nxall=nrow(xall)  
  xall=as.matrix(xall[1:(nxall-1),])  
  allpress=apply(xall,1,getpress,y=y,Xmat=X)  
  list(which=xall,press=allpress)  
}
```

PRESS Statistic: Wine Example

```
> X=wines[,-12]
> pressmod=presslm(wines$quality,X)
> widx=which.min(pressmod$press)
> xidx=(1:ncol(X))[pressmod$which[widx,]]
> Xin=data.frame(X[,xidx])
> pressmod=lm(wines$quality~.,data=Xin)
> summary(pressmod)
```

Call:

```
lm(formula = wines$quality ~ ., data = Xin)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.68918	-0.36757	-0.04653	0.46081	2.02954

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.4300987	0.4029168	10.995	< 2e-16 ***
volatile.acidity	-1.0127527	0.1008429	-10.043	< 2e-16 ***
chlorides	-2.0178138	0.3975417	-5.076	4.31e-07 ***
free.sulfur.dioxide	0.0050774	0.0021255	2.389	0.017 *
total.sulfur.dioxide	-0.0034822	0.0006868	-5.070	4.43e-07 ***
pH	-0.4826614	0.1175581	-4.106	4.23e-05 ***
sulphates	0.8826651	0.1099084	8.031	1.86e-15 ***
alcohol	0.2893028	0.0167958	17.225	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6477 on 1591 degrees of freedom

Multiple R-squared: 0.3595, Adjusted R-squared: 0.3567

F-statistic: 127.6 on 7 and 1591 DF, p-value: < 2.2e-16

R State Facts Data

The `state.x77` matrix contains 8 variables (columns) collected from the 50 states (rows) during the early 1970s

- `Population`: estimate of state population (1975)
- `Income`: per capita income (1974)
- `Illiteracy`: percent illiterate (1970)
- `Life Exp`: life expectancy (1969–1971)
- `Murder`: murder rate per 100,000 people (1976)
- `HS Grad`: percent high-school graduates (1970)
- `Frost`: mean number of days with minimum temperature below freezing (1931–1960)
- `Area`: land area in square miles

Look at State Facts Data

```
> states=data.frame(state.x77,row.names=state.abb)
> states[1:15,]
```

	Population	Income	Illiteracy	Life.Exp	Murder	HS.Grad	Frost	Area
AL	3615	3624	2.1	69.05	15.1	41.3	20	50708
AK	365	6315	1.5	69.31	11.3	66.7	152	566432
AZ	2212	4530	1.8	70.55	7.8	58.1	15	113417
AR	2110	3378	1.9	70.66	10.1	39.9	65	51945
CA	21198	5114	1.1	71.71	10.3	62.6	20	156361
CO	2541	4884	0.7	72.06	6.8	63.9	166	103766
CT	3100	5348	1.1	72.48	3.1	56.0	139	4862
DE	579	4809	0.9	70.06	6.2	54.6	103	1982
FL	8277	4815	1.3	70.66	10.7	52.6	11	54090
GA	4931	4091	2.0	68.54	13.9	40.6	60	58073
HI	868	4963	1.9	73.60	6.2	61.9	0	6425
ID	813	4119	0.6	71.87	5.3	59.5	126	82677
IL	11197	5107	0.9	70.14	10.3	52.6	127	55748
IN	5313	4458	0.7	70.88	7.1	52.9	122	36097
IA	2861	4628	0.5	72.56	2.3	59.0	140	55941

State Data: Backward Elimination

```
> bsmod=bslm(states$Murder,states[,-5])
[1] "Step 1 : Dropping Income"
[1] "Step 2 : Dropping HS.Grad"
> summary(bsmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.202e+02	1.718e+01	6.994	1.17e-08	***
Population	1.780e-04	5.930e-05	3.001	0.00442	**
Illiteracy	1.173e+00	6.801e-01	1.725	0.09161	.
Life.Exp	-1.608e+00	2.324e-01	-6.919	1.50e-08	***
Frost	-1.373e-02	7.080e-03	-1.939	0.05888	.
Area	6.804e-06	2.919e-06	2.331	0.02439	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.712 on 44 degrees of freedom
 Multiple R-squared: 0.8068, Adjusted R-squared: 0.7848
 F-statistic: 36.74 on 5 and 44 DF, p-value: 1.221e-14

State Data: Forward Selection

```
> fsmod=fslm(states$Murder,states[,-5])
[1] "Step 1 : Adding Life.Exp"
[1] "Step 2 : Adding Frost"
[1] "Step 3 : Adding Population"
[1] "Step 4 : Adding Area"
[1] "Step 5 : Adding Illiteracy"
> summary(fsmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.202e+02	1.718e+01	6.994	1.17e-08	***
Life.Exp	-1.608e+00	2.324e-01	-6.919	1.50e-08	***
Frost	-1.373e-02	7.080e-03	-1.939	0.05888	.
Population	1.780e-04	5.930e-05	3.001	0.00442	**
Area	6.804e-06	2.919e-06	2.331	0.02439	*
Illiteracy	1.173e+00	6.801e-01	1.725	0.09161	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.712 on 44 degrees of freedom
 Multiple R-squared: 0.8068, Adjusted R-squared: 0.7848
 F-statistic: 36.74 on 5 and 44 DF, p-value: 1.221e-14

State Data: Stepwise Selection

```
> swmod=swlm(states$Murder,states[, -5])
[1] "Step 1 : Adding Life.Exp"
[1] "Step 2 : Adding Frost"
[1] "Step 3 : Adding Population"
[1] "Step 4 : Adding Area"
[1] "Step 5 : Adding Illiteracy"
> summary(swmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.202e+02	1.718e+01	6.994	1.17e-08	***
Life.Exp	-1.608e+00	2.324e-01	-6.919	1.50e-08	***
Frost	-1.373e-02	7.080e-03	-1.939	0.05888	.
Population	1.780e-04	5.930e-05	3.001	0.00442	**
Area	6.804e-06	2.919e-06	2.331	0.02439	*
Illiteracy	1.173e+00	6.801e-01	1.725	0.09161	.

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.712 on 44 degrees of freedom
 Multiple R-squared: 0.8068, Adjusted R-squared: 0.7848
 F-statistic: 36.74 on 5 and 44 DF, p-value: 1.221e-14

State Data: Adjusted R^2 Selection

```
> X=states[,-5]
> arsqmod=leaps(x=X,y=states$Murder,method="adjr2")
> widx=which.max(arsqmod$adjr2)
> xidx=(1:ncol(X))[arsqmod$which[widx,]]
> Xin=data.frame(X[,xidx])
> arsqmod=lm(states$Murder~.,data=Xin)
> summary(arsqmod)          # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.202e+02	1.718e+01	6.994	1.17e-08	***
Population	1.780e-04	5.930e-05	3.001	0.00442	**
Illiteracy	1.173e+00	6.801e-01	1.725	0.09161	.
Life.Exp	-1.608e+00	2.324e-01	-6.919	1.50e-08	***
Frost	-1.373e-02	7.080e-03	-1.939	0.05888	.
Area	6.804e-06	2.919e-06	2.331	0.02439	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.712 on 44 degrees of freedom
 Multiple R-squared: 0.8068, Adjusted R-squared: 0.7848
 F-statistic: 36.74 on 5 and 44 DF, p-value: 1.221e-14

State Data: Stepwise AIC Selection

```
> smod=lm(states$Murder~.,data=states)
> aicmod=step(smod,trace=0)
> summary(aicmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.202e+02	1.718e+01	6.994	1.17e-08	***
Population	1.780e-04	5.930e-05	3.001	0.00442	**
Illiteracy	1.173e+00	6.801e-01	1.725	0.09161	.
Life.Exp	-1.608e+00	2.324e-01	-6.919	1.50e-08	***
Frost	-1.373e-02	7.080e-03	-1.939	0.05888	.
Area	6.804e-06	2.919e-06	2.331	0.02439	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.712 on 44 degrees of freedom

Multiple R-squared: 0.8068, Adjusted R-squared: 0.7848

F-statistic: 36.74 on 5 and 44 DF, p-value: 1.221e-14

State Data: Stepwise BIC Selection

```
> smod=lm(states$Murder~.,data=states)
> bicmod=step(smod,k=log(50),trace=0)
> summary(bicmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.387e+02	1.369e+01	10.136	3.40e-13	***
Population	1.581e-04	5.944e-05	2.660	0.010778	*
Life.Exp	-1.837e+00	1.946e-01	-9.442	3.04e-12	***
Frost	-2.204e-02	5.299e-03	-4.160	0.000141	***
Area	7.387e-06	2.962e-06	2.494	0.016374	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.75 on 45 degrees of freedom

Multiple R-squared: 0.7937, Adjusted R-squared: 0.7754

F-statistic: 43.28 on 4 and 45 DF, p-value: 7.106e-15

State Data: Mallow's C_p Selection

```
> X=states[,-5]
> cpmod=leaps(x=X,y=states$Murder,method="Cp")
> widx=which.min(cpmod$Cp)
> xidx=(1:ncol(X))[cpmod$which[widx,]]
> Xin=data.frame(X[,xidx])
> cpmod=lm(states$Murder~.,data=Xin)
> summary(cpmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.202e+02	1.718e+01	6.994	1.17e-08	***
Population	1.780e-04	5.930e-05	3.001	0.00442	**
Illiteracy	1.173e+00	6.801e-01	1.725	0.09161	.
Life.Exp	-1.608e+00	2.324e-01	-6.919	1.50e-08	***
Frost	-1.373e-02	7.080e-03	-1.939	0.05888	.
Area	6.804e-06	2.919e-06	2.331	0.02439	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.712 on 44 degrees of freedom
 Multiple R-squared: 0.8068, Adjusted R-squared: 0.7848
 F-statistic: 36.74 on 5 and 44 DF, p-value: 1.221e-14

State Data: PRESS Selection

```
> X=states[,-5]
> pressmod=presslm(states$Murder,X)
> widx=which.min(pressmod$press)
> xidx=(1:ncol(X)) [pressmod$which[widx,]]
> Xin=data.frame(X[,xidx])
> pressmod=lm(states$Murder~.,data=Xin)
> summary(pressmod)      # I deleted some output
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.259e+02	1.777e+01	7.083	8.64e-09	***
Population	1.946e-04	6.078e-05	3.202	0.00254	**
Illiteracy	1.912e+00	7.620e-01	2.509	0.01587	*
Life.Exp	-1.757e+00	2.491e-01	-7.053	9.57e-09	***
HS.Grad	7.626e-02	4.369e-02	1.746	0.08786	.
Frost	-1.011e-02	7.199e-03	-1.404	0.16719	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.755 on 44 degrees of freedom

Multiple R-squared: 0.797, Adjusted R-squared: 0.7739

F-statistic: 34.54 on 5 and 44 DF, p-value: 3.565e-14